

IMPLEMENTING NONLINEAR OPTIMIZATION ALGORITHMS IN PYTHON

Truncated Newton · Sequential Penalty · Filled Functions

Continuous Optimization · Alessandro Rapo · Sapienza University of Rome · 2024–2025

01 — ABSTRACT

This article documents the Python implementation of three fundamental algorithms from continuous nonlinear optimization: the Truncated Newton method, the Sequential Penalty method, and the Filled Function algorithm for global optimization. Every component — from the inner conjugate gradient solver to the Armijo line search and the Gaussian hill construction — was built from scratch in NumPy, with particular attention to numerical stability and adherence to theoretical convergence guarantees.

02 — INTRODUCTION

Nonlinear optimization is at the heart of machine learning, engineering design, financial modeling, and operations research. While production-grade solvers such as IPOPT, KNITRO, or `scipy.optimize` are widely available, implementing the underlying algorithms by hand remains one of the most effective ways to develop a deep, operational understanding of their behavior — including their failure modes and sensitivity to parameter choices.

The project addresses three distinct but related problems:

- **NT:** Unconstrained local minimization via the Truncated Newton method.
- **PS:** Constrained local minimization via Sequential Penalty functions.
- **FF:** Global minimization via additive Filled Functions.

All three algorithms are implemented in pure NumPy without automatic differentiation. Gradients are provided analytically or approximated with central finite differences, and Hessian-vector products are computed via a matrix-free finite-difference scheme, avoiding the $O(n^2)$ cost of assembling the full Hessian.

03 — TRUNCATED NEWTON METHOD (NEWTONT.PY)

The Truncated Newton method is designed for large-scale unconstrained optimization. At each outer iteration k , instead of solving the Newton system $H_k d = -g_k$ exactly, an approximate solution is computed using the Conjugate Gradient method (Steihaug-CG), truncating early when negative curvature is detected or when the residual is small enough.

► ***Hessian-Vector Products***

The key computational primitive is the matrix-free approximation of $H(x) \cdot v$ via central finite differences on the gradient:

$$H(x) \cdot v \approx [\nabla f(x + h \cdot v^\wedge) - \nabla f(x - h \cdot v^\wedge)] / (2h / \|v\|) \quad h = 1e-5$$

This avoids forming the $n \times n$ Hessian explicitly and keeps memory usage at $O(n)$, which is critical for high-dimensional problems.

► ***Truncated Conjugate Gradient (GCT)***

The inner CG loop solves $H \cdot d \approx -g$ starting from $d = 0$. Two truncation criteria are enforced:

- Negative curvature: if $s^\wedge T H s \leq 0$, return the current iterate (descent is guaranteed).
- Relative residual: stop when $\|r_k\| \leq \eta_k \cdot \|g\|$, with $\eta_k = \min(0.5, \sqrt{\|g\|})$, ensuring superlinear convergence.

► ***Armijo Line Search***

$$f(x + \alpha \cdot d) \leq f(x) + \alpha \cdot \gamma \cdot \nabla f(x)^\wedge T d \quad \gamma = 1e-4, \quad \beta = 0.5$$

If the CG direction is not a descent direction (possible near saddle points), the algorithm falls back to the steepest descent direction $-g$.

04 — SEQUENTIAL PENALTY METHOD (PENALITA.PY)

The Sequential Penalty method converts a constrained problem into a sequence of unconstrained ones by penalizing constraint violations. Given $\min f(x)$ s.t. $g(x) \leq 0$, $h(x) = 0$, the penalized objective is:

$$F_\varepsilon(x) = f(x) + (1/\varepsilon) \cdot \sum \max(0, g_i(x))^2 + (1/\varepsilon) \cdot \sum h_j(x)^2$$

As $\varepsilon \rightarrow 0$, the penalty weight $1/\varepsilon \rightarrow \infty$, forcing the minimizer of F_ε towards the feasible region. The algorithm follows the four-step schema from the lecture notes:

- P1: Estimate KKT multipliers as $\lambda_i = (2/\varepsilon) \cdot \max(0, g_i(x))$ and $\mu_j = (2/\varepsilon) \cdot h_j(x)$. Check KKT conditions.
- P2: Minimize F_ε approximately using Truncated Newton, stopping when $\|\nabla F_\varepsilon\| \leq \delta$.
- P3: If constraint violation did not decrease by factor θ_1 , reduce $\varepsilon \leftarrow \theta_2 \cdot \varepsilon$ (increase penalty weight).
- P4: Tighten inner tolerance: $\delta \leftarrow \theta_3 \cdot \delta$. Repeat from P1.

► **KKT Verification**

KKT conditions are checked across four criteria at each outer iteration: stationarity of the Lagrangian ($\|\nabla L\| \leq \text{tol}$), primal feasibility, dual feasibility ($\lambda \geq 0$), and complementary slackness ($\lambda_i \cdot g_i(x) = 0$). Termination occurs as soon as all four hold within tolerance.

► **Parameter Choices**

We use $\varepsilon_0 = 1.0$ (moderate initial penalty), $\theta_1 = 0.25$ (strict improvement threshold), $\theta_2 = 0.1$ (aggressive penalty growth), and $\theta_3 = 0.5$ (gradual tightening). This aggressive schedule avoids the ill-conditioning that arises when ε is reduced too slowly.

05 – FILLED FUNCTION ALGORITHM (FILLED.PY)

Local optimization methods converge to the nearest local minimum, which may be far from the global one. The Filled Function method provides a principled escape strategy: construct an auxiliary function that eliminates the current local minimum and guides the search towards a better basin.

► **The Additive Filled Function – Type 1**

Given the current local minimizer x^*_k , the function is:

$$U_k(x) = \tau \cdot (\min\{0, f(x) - f(x^*_k) + \rho\})^3 + \exp(-\|x - x^*_k\|^2 / \gamma^2)$$

The two terms serve complementary roles:

- Gaussian term: creates a hill centered at x^*_k , making it a strict local maximum of U_k – any local minimizer of U_k starting from x^*_k is forced to move elsewhere.
- Cubic penalty: is zero where $f(x) \geq f(x^*_k) - \rho$ and strongly negative where $f(x) < f(x^*_k) - \rho$, pulling the search towards regions of lower objective value.

For sufficiently large τ and $0 < \rho < f(x^*_k) - f^*$, all global minimizers of U_k lie in $\{x : f(x) < f(x^*_k)\}$, guaranteeing progress towards the global optimum.

► **Multi-Start Strategy**

- Geometric perturbations along each coordinate axis at scales $\gamma \cdot 10^0, \gamma \cdot 10^{-1}, \gamma \cdot 10^{-2}, \gamma \cdot 10^{-3}$.
- Uniformly random restarts sampled across the feasible box $[box_lo, box_hi]^n$.
- Any candidate z with $f(z) < f(x^*_k)$ is accepted, then refined with a full NT run on f .

06 — EXPERIMENTAL RESULTS

► **Sequential Penalty — Selected Problems**

- **P3 (Rosenbrock + constraints):** $x^* = (0.5, 0.866)$, $f^* = 38.199$. Active constraints: circle $g_3 = 0$ and bound $g_4 = 0$.
- **P7 (minimum distance + equalities):** KKT at iteration 10. $x^* \approx (2.0, 2.0, 0.849, 1.131)$, $f^* = 13.858$.
- **P11 (maximize x_1 + mixed):** KKT at iteration 15. $x^* = (1, 1, 0, 0)$, $f^* = -1.0$ — global optimum confirmed.

► **Filled Functions — Selected Problems**

- **F1 (Rosenbrock 4D):** $f^* = 0.0$ at $x^* = (1,1,1,1)$ — global minimum confirmed.
- **F5 (Wood function):** $f^* = 0.0$ at $x^* = (1,1,1,1)$ — global minimum confirmed.
- **F6 (Six-Hump Camel):** $f^* = -1.031628$ at $x^* \approx (-0.090, 0.713)$ — matches known global minimum.
- **F129 (Box volume):** $f^* = -3456.0$ at $x^* = (12, 12)$ — analytic global minimum recovered.

07 — IMPLEMENTATION NOTES & LESSONS LEARNED

- **Finite difference step size:** $h = 1e-5$ for Hessian-vector products and gradients. Too small ($< 1e-8$) causes cancellation; too large ($> 1e-3$) corrupts the CG direction.
- **Adaptive CG tolerance:** $\eta_k = \min(0.5, \sqrt{\|g_k\|})$ rather than fixed — coarse solutions early, refined solutions near optimality — dramatically improves convergence.
- **Filled function γ sensitivity:** Controls the basin of repulsion around x^*_k . Problem-specific tuning ($\gamma \in \{0.05, 0.5, 1.0, 2.0\}$) proved necessary — no single value works universally.

- **Penalty ill-conditioning:** As $\varepsilon \rightarrow 0$ the Hessian of F_ε becomes ill-conditioned. Adaptive inner tolerance $\delta \leftarrow \theta_3 \cdot \delta$ compensates by requiring less precision when the landscape is steep.

08 — CONCLUSION

Implementing these algorithms from scratch — without auto-diff or external solvers — offers a level of insight that using a library simply cannot provide. Every hyperparameter has a theoretical justification rooted in convergence proofs; every failure mode reveals something concrete about the geometry of the problem.

The Truncated Newton inner solver proved to be the critical shared component across all three algorithms: a robust, numerically stable GCT implementation is what enables both the penalty method and the filled function framework to converge reliably. The full implementation — `funzioni.py`, `NewtonT.py`, `penalita.py`, `Filled.py` — covers 20+ benchmark problems across constrained and global optimization.

— — REFERENCES

- [1] S. Lucidi, Appunti dalle lezioni di Ottimizzazione Continua, Sapienza, A.A. 2024–2025.
- [2] J. Nocedal, S. J. Wright, Numerical Optimization, 2nd ed., Springer, 2006.
- [3] T. Steihaug, The conjugate gradient method and trust regions in large scale optimization, SIAM J. Numer. Anal., 1983.
- [4] Y.-P. Zhang, Filled functions for unconstrained global optimization, J. Global Optim., 2001.
- [5] R. Fletcher, Practical Methods of Optimization, 2nd ed., Wiley, 1987.